

Basics of Kubernetes



Agenda:

- Introduction
- Containers Orchestration Tools
- Kubernetes origin and definition
- Kubernetes Architecture
- Kubernetes Server Roles
- Kubernetes Cluster
- Kubernetes Components
- Kubernetes Basic Objects
- Benefits of Kubernetes
- Kubectl Definition and Basic Commands

Introduction

● Immutable Infrastructure

- Never modifies directly.
- For making any change to server:
 - Create copy of server from base image -> changes made -> replace old with new one

● Containers

- Offers a way to package code, runtime, system tools, system libraries, and configs altogether.
- Lightweight shipment, standalone executable.
- Consistent application behaviour on multiple environments

Containers Orchestrators



Docker Swarm



Kubernetes



MESOS

Kubernetes



Origin:

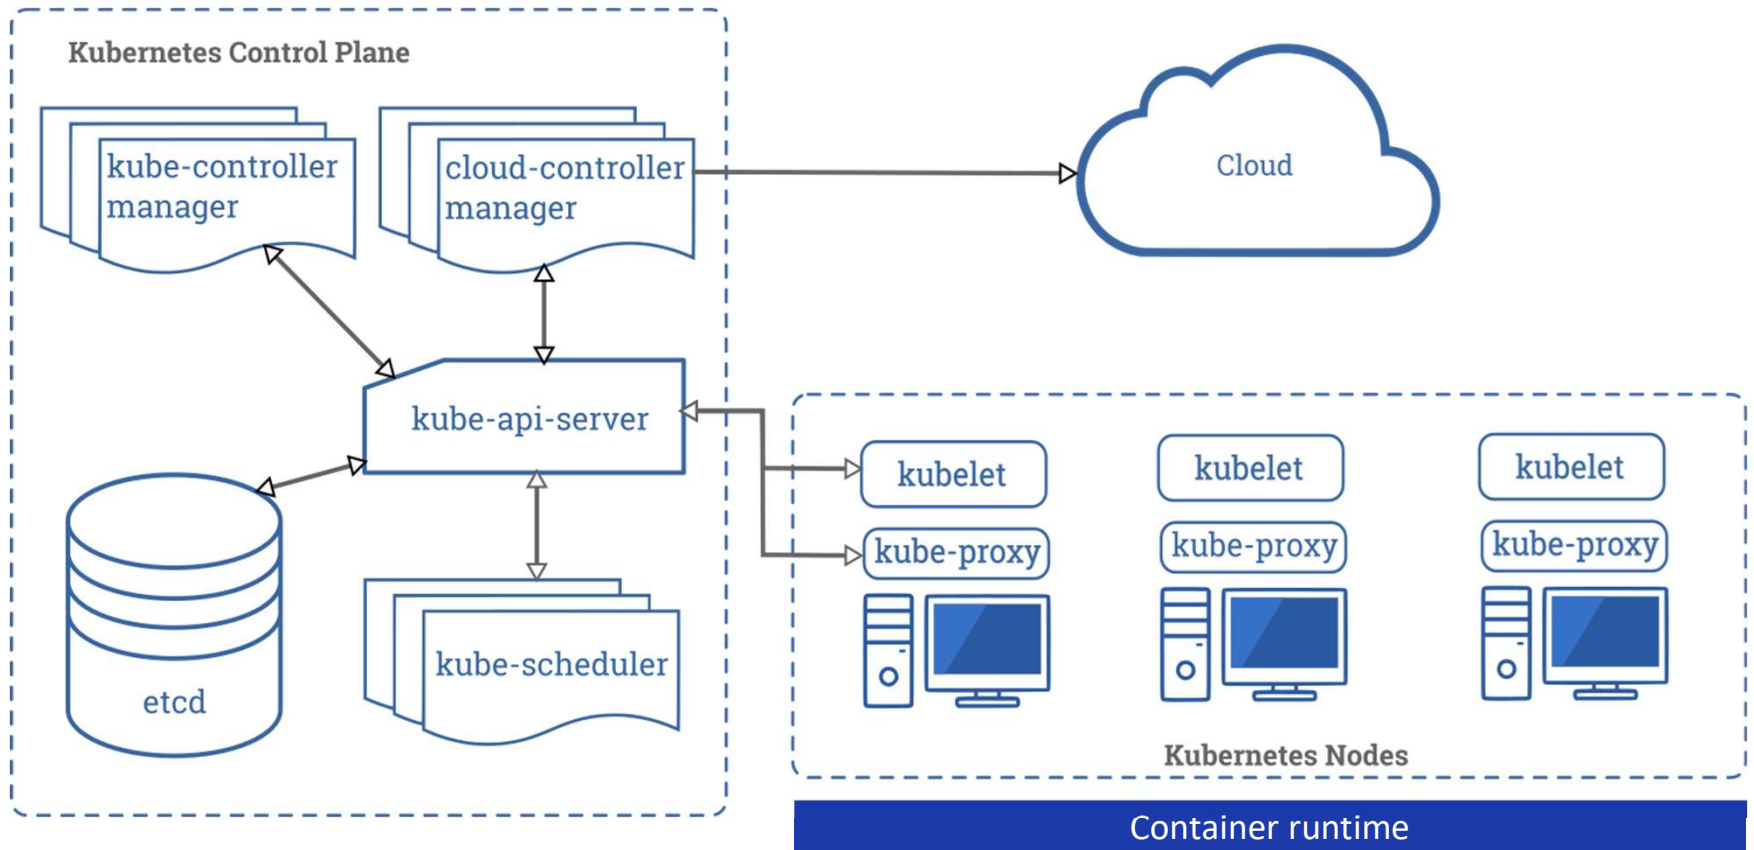
- Greek word - means **helmsman or pilot**
- Abbreviated as **k8s**
- Open sourced by Google in 2014

Definition:

“Kubernetes is a portable, extensible, open-source platform for managing containerized workloads and services, that facilitates both declarative configuration and automation”

<https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/>

Kubernetes Architecture



Control Plane Components

The control plane's components make global decisions about the cluster (for example, scheduling), as well as detecting and responding to cluster events (for example, starting up a new pod when a deployment's replicas field is unsatisfied).

Control plane components can be run on any machine in the cluster. However, for simplicity, set up scripts typically start all control plane components on the same machine, and do not run user containers on this machine.

kube-controller-manager

Think of the `kube-controller-manager` in Kubernetes as the caretaker that watches over your applications and keeps them in good shape. It makes sure everything is working correctly, like a gardener who tends to a garden, ensuring the plants grow well, stay healthy, and are in the right place. The controller manager does a similar job, tending to your applications, making sure they're running smoothly and according to your desired settings and rules.

Cloud-controller-manager

The `cloud-controller-manager` is like a helpful assistant that takes care of all the special tasks needed to make Kubernetes work smoothly with your cloud services, so you don't have to worry about the technical details. It's like having someone who speaks both Kubernetes and your cloud provider's language, ensuring everything works well together.

Kube-apiserver

- The `kube-apiserver` is like the front door of your Kubernetes cluster. It's the component that acts as a gateway for communication between users, applications, and the cluster itself. It receives requests and commands and makes sure they are processed correctly within the cluster. Think of it as the security guard at a club entrance, ensuring only the right people get in and that they follow the rules.

Kube-scheduler

- The `kube-scheduler` is a Kubernetes control plane component responsible for selecting an appropriate node (worker machine) in the cluster for running a new pod (a unit of work in Kubernetes). It does this by considering factors like resource requirements, affinity/anti-affinity rules, and other policies defined in the scheduling process. The goal is to efficiently distribute workloads across the cluster while satisfying constraints and requirements, ensuring optimal resource utilization and application performance.

etcd

- Consistent and highly-available key value store used as Kubernetes' backing store for all cluster data.
- If your Kubernetes cluster uses etcd as its backing store, make sure you have a [back up](#) plan for those data.

kubelet

- An agent that runs on each [node](#) in the cluster. It makes sure that [containers](#) are running in a [Pod](#).
- The kubelet takes a set of PodSpecs that are provided through various mechanisms and ensures that the containers described in those PodSpecs are running and healthy. The kubelet doesn't manage containers which were not created by Kubernetes.

kube-proxy

- kube-proxy is a network proxy that runs on each [node](#) in your cluster, implementing part of the Kubernetes [Service](#) concept.
- [kube-proxy](#) maintains network rules on nodes. These network rules allow network communication to your Pods from network sessions inside or outside of your cluster.
- kube-proxy uses the operating system packet filtering layer if there is one and it's available. Otherwise, kube-proxy forwards the traffic itself.

Container runtime

- The container runtime is the software that is responsible for running containers.
- Kubernetes supports container runtimes such as [containerd](#), [CRI-O](#), and any other implementation of the [Kubernetes CRI \(Container Runtime Interface\)](#).

Kubernetes Server Roles

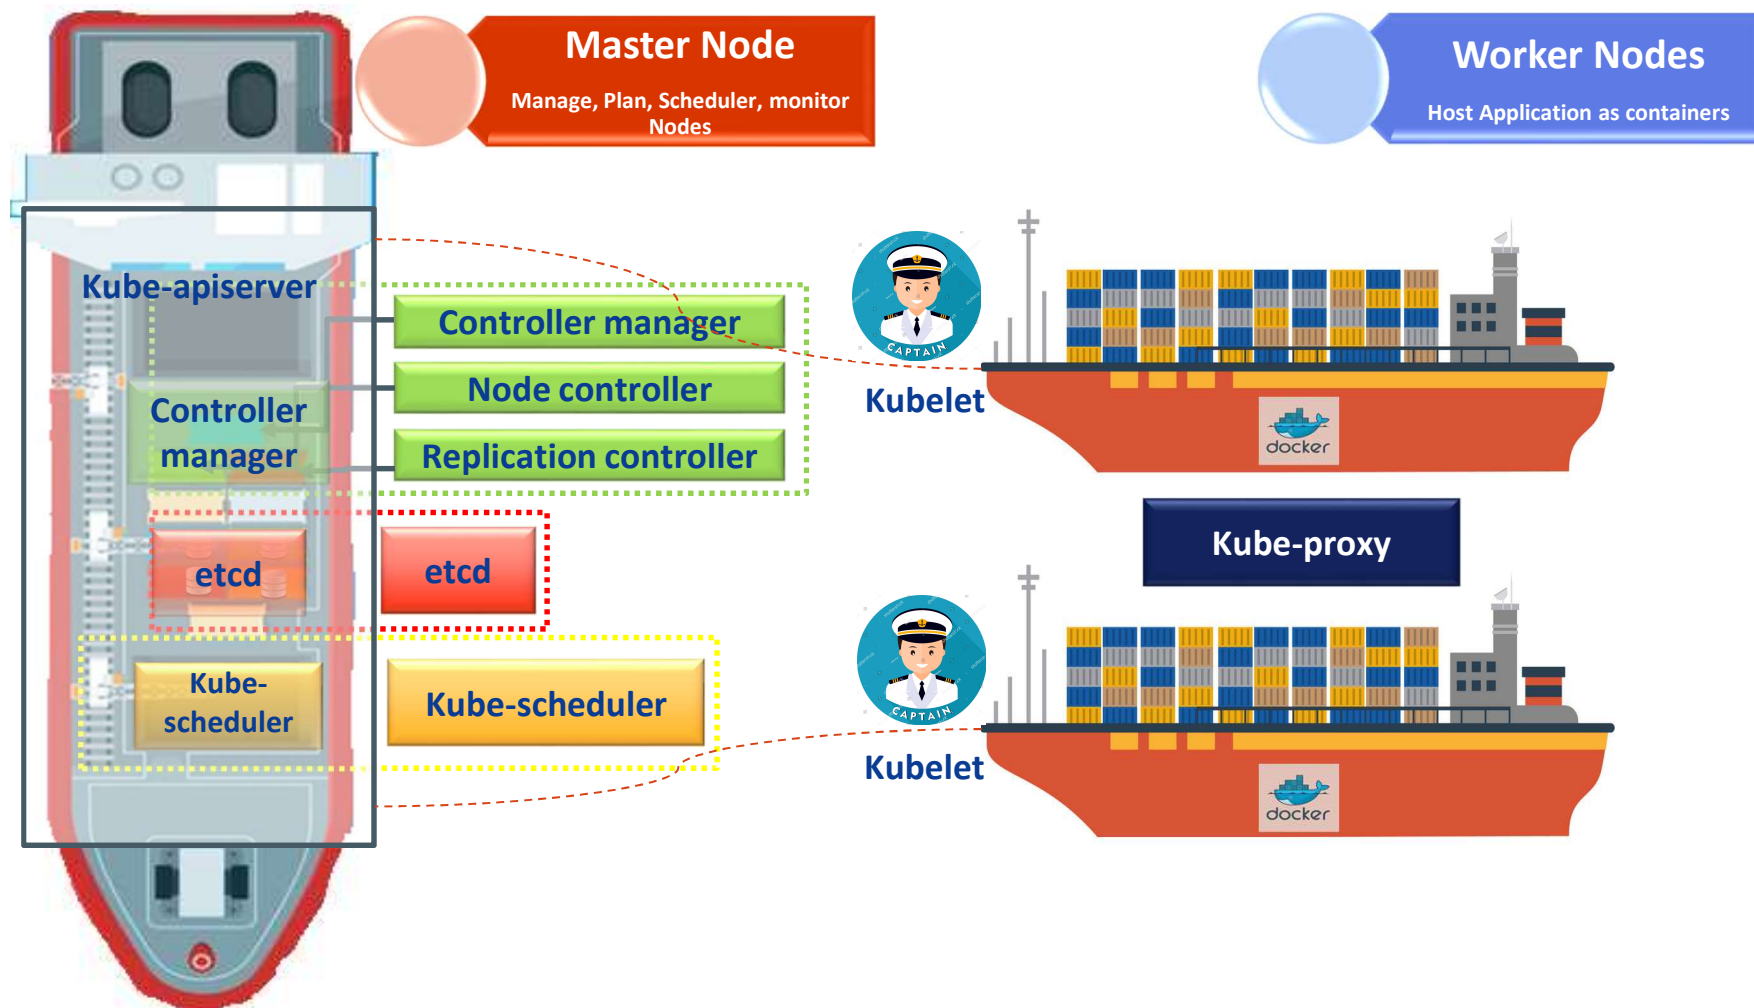
- **Master**

- Gateway and brain for the cluster
- Exposes an API for users and clients
- Responsible for health checking , scheduling
- Orchestration communication between other components

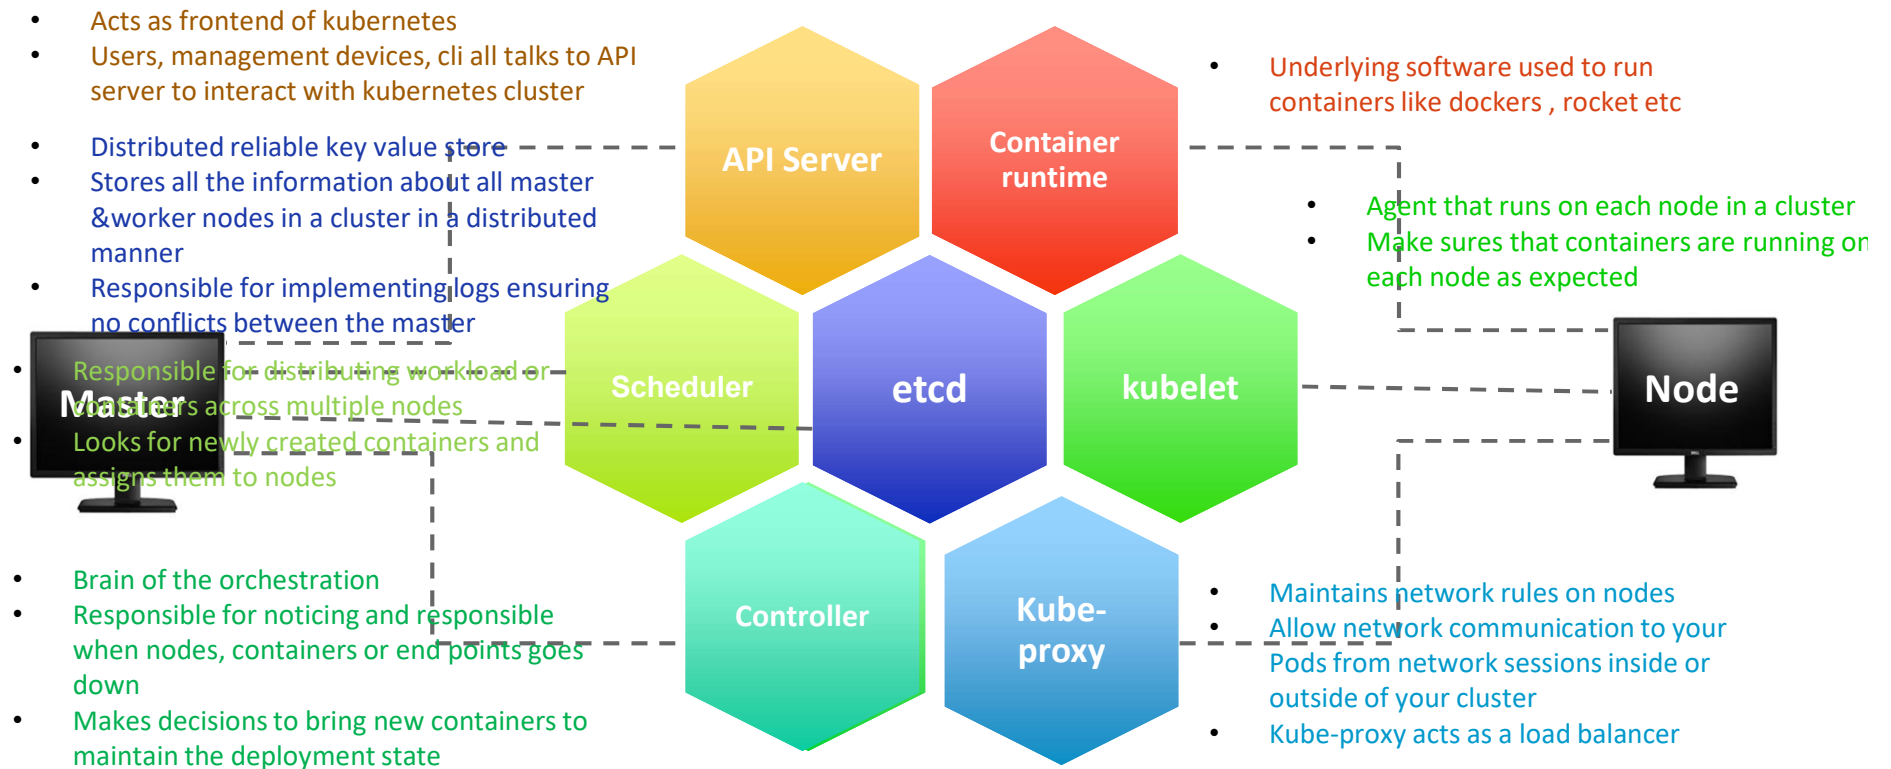
- **Node**

- Accepts and runs workloads
- Receives work instructions from the master server and creates or destroys containers accordingly
- Adjusts networking rules to route and forward traffic appropriately

Kubernetes Cluster



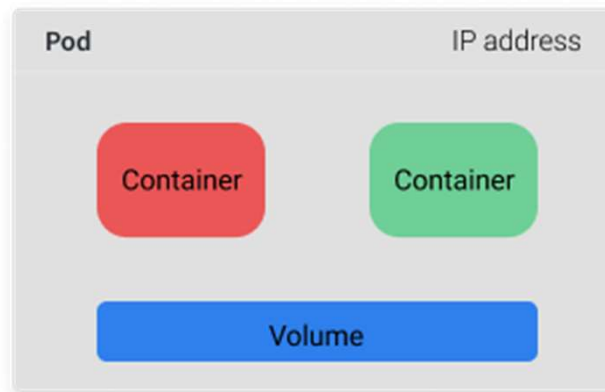
Kubernetes Components



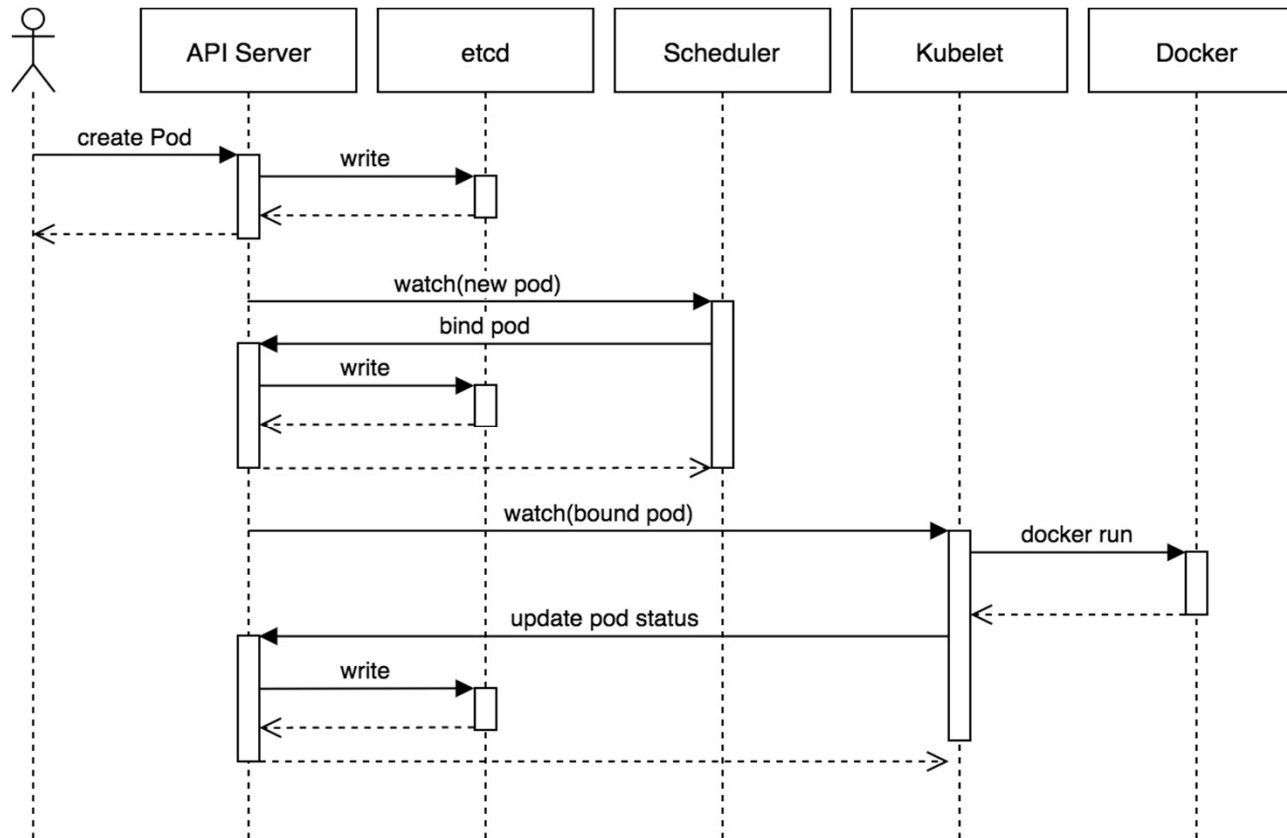
Kubernetes Basic Objects

- **Pod**

- Fundamental building block in Kubernetes
- Comprised of one or more (tightly related) containers, a shared networking layer, and shared file system volumes
- Similar to containers, pods are designed to be ephemeral

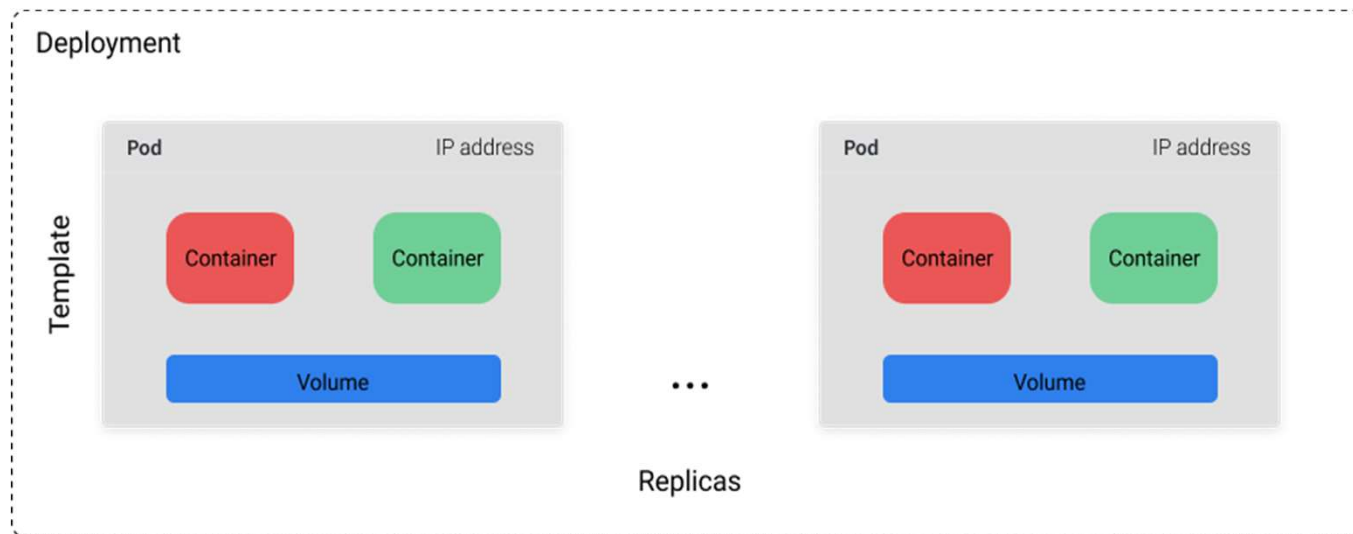


The following sequence diagram, from [Heptio's](#) blog, shows the components involved during a simple Pod creation process. It's a great illustration of the API Server and etcd interaction.



- **Deployment**

- Contains a collection of pods defined by a template and a replica
- Best suited for stateless applications
- Allow us to specify the strategy of rolling updates when we have new versions of our container image



Deployment Example

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: ml-model-serving
  labels:
    app: ml-model
spec:
  replicas: 10
  selector:
    matchLabels:
      app: ml-model
  template:
    metadata:
      labels:
        app: ml-model
    spec:
      containers:
        - name: ml-rest-server
          image: ml-serving:1.0
          ports:
            - containerPort: 80
```

How many Pods should be running?

How do we find Pods that belong to this Deployment?

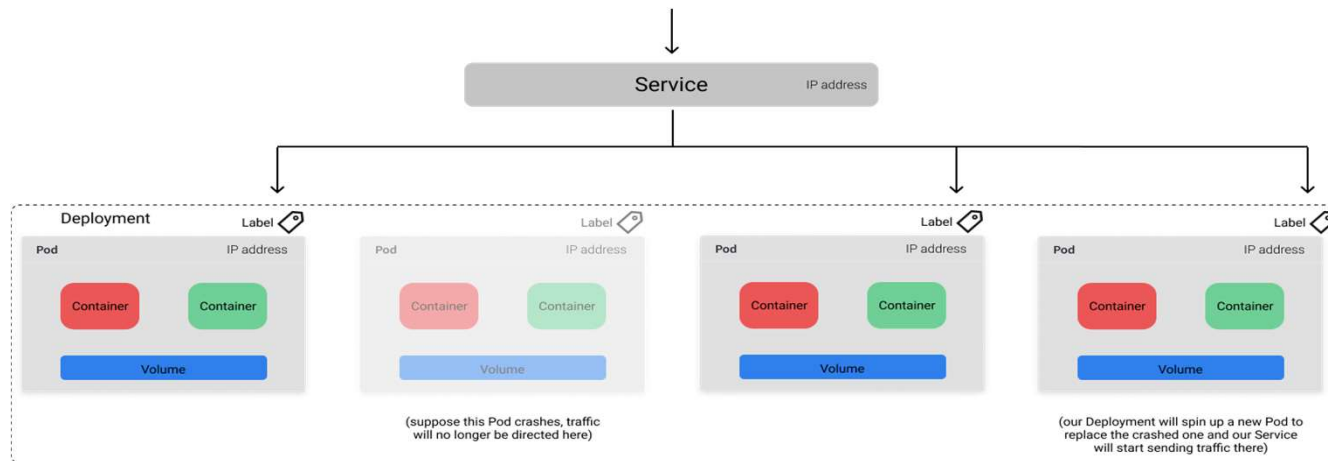
What should a Pod look like?

Add a label to the Pods so our Deployment can find the Pods to manage.

What containers should be running in the Pod?

- **Service**

- Provides with a stable endpoint which can be used to direct traffic to the desired Pods even as the exact underlying Pods change due to updates, scaling, and failures
- Services know which Pods they should send traffic to based on labels (key-value pairs) which is defined in the Pod metadata.



In this example, Service sends traffic to all healthy Pods with the label `app="ml-model"`

Service Example



```
apiVersion: v1
kind: Service
metadata:
  name: ml-model-svc
  labels:
    app: ml-model
spec:
  type: ClusterIP
  selector:
    app: ml-model
  ports:
    - protocol: TCP
      port: 80
```

How do we want to expose our endpoint?

How do we find Pods to direct traffic to?

How will clients talk to our Service?

Some other objects , include:

- **ReplicaSet (RS)**. Ensures the **desired amount of pod** is what's running.
- **StatefulSet**. A workload API object that manages stateful applications, such as databases.
- **DaemonSet**. Ensures that all or some worker nodes run a copy of a pod. This is useful for daemon applications like **Fluentd**.
- **Job**. Creates one or more pods, runs a certain task(s) to completion, then deletes the pod(s).
- **Volume**. An abstraction that lets us persist data. (This is necessary because containers are ephemeral—meaning data is deleted when the container is deleted.)
- **Namespace**. A segment of the cluster dedicated to a certain purpose, for example a certain project or team of devs.

Benefits of Kubernetes

- **Horizontal scaling.** Scale your application [as needed](#) from command line or UI.
- **Automated rollouts and rollbacks.** Roll out changes that monitor the health of your application—ensuring all instances don't fail or go down simultaneously. If something goes wrong, K8S automatically rolls back the change.
- **Service discovery and load balancing.** Containers get their own IP so you can put a set of containers behind a single DNS name for load balancing.
- **Storage orchestration.** Automatically mount local or public cloud or a network storage.
- **Secret and configuration management.** [Create and update secrets](#) and configs without rebuilding your image.
- **Self-healing.** The platform heals many problems: restarting failed containers, replacing and rescheduling containers as nodes die, killing containers that don't respond to your user-defined health check, and waiting to advertise containers to clients until they're ready.
- **Automatic binpacking.** Automatically schedules containers based on resource requirements and other constraints.

Kubectl

Kube control tool/Kube command line tool

- Kubernetes CLI tool
- Used to deploy and manage applications on a kubernetes cluster
- Gets cluster related information
- Gets status of the nodes in the cluster and many others

Kubectl Basic Commands

- `kubectl get` - list resources
- `kubectl describe` - show detailed information about a resource
- `kubectl logs` - print the logs from a container in a pod
- `kubectl exec` - execute a command on a container in a pod
- `Kubectl run hello-minikube` - Deploys application in a cluster
- `Kubectl cluster-info` - View information about the cluster
- `Kubectl get nodes` - View all the nodes part of a cluster
- `Kubectl get deployments` - View information about the deployments

Thank you